

R Programming: Worksheet 6

1. Straightforward examples

For the following, perform the below in a parallel manner, using 2 or more cores. Here you can choose between either a fork approach using `mclapply` or similar, or a socket approach using `makeCluster` and `parSapply` or similar. If you're on a Windows computer you'll need to use the sockets approach for true parallel computing.

- (a) For 100 replicates each, calculate the average of 10000 draws from a random distribution with mean 0 and standard deviation 1
- (b) Calculate the column average for each column of the built in dataset `mtcars`

2. Bagging

Bootstrap aggregating, or bagging, is a technique in machine learning designed to improve the accuracy of predictors, as well as reduce the variance in the estimated predictions and help to reduce overfitting. In brief, semi-formally, given some training set D of size n , bagging generates m new training sets D_i , by sampling from D uniformly with replacement. Then, m models are fitted, one on each D_i , and prediction is done using averaging over the predictions from the m models.

We use data from <https://www.kaggle.com/uciml/autompg-dataset> that has been lightly modified that you can download as `autompg_clean.csv` from the course website. We will try to predict `mpg` (miles per gallon) using the other variables, using linear regression, and bagging. Here we will use parallel computing to speed up the model building step.

- (a) Read in `autompg_dataset.csv`, and have a quick look at the contents. Sample 80% of the rows without replacement to form a training dataset, and make a test dataset with the remaining samples *Note that `read.csv(file)` or similar will produce a column called `X`, which is the row number. You can either use `read.csv(file, row.names = 1)` to load it in, or just generally discard it after loading in the data.*
- (b) Write a function that performs linear regression using `lm` using a bootstrapped sample of the training data. Predict `mpg` using all other variables, except `car.name`
- (c) Using parallel computing, run this function 100 times, to generate 100 models using the 100 bootstrapped training samples
- (d) *Optional* Build a function that generates the bagged estimate on the test set data, and compare the correlation between the bagged model and the truth `test[, "mpg"]`. Compare this to a doing a more conventional single linear regression using all the training data once.

3. Bootstrap hypothesis testing

Here we will consider the following algorithm for comparing the means of two samples (from Efron, B.; Tibshirani, R. (1993). An Introduction to the Bootstrap.)

See also [https://en.wikipedia.org/wiki/Bootstrapping_\(statistics\)](https://en.wikipedia.org/wiki/Bootstrapping_(statistics)) the Wikipedia entry for bootstrapping.

Let x_1, \dots, x_n be a random sample from some distribution F with sample mean \bar{x} and sample variance σ_x^2 . Let y_1, \dots, y_m be another, independent random sample from some distribution G with mean \bar{y} and sample variance σ_y^2 .

i Calculate

$$t = \frac{\bar{x} - \bar{y}}{\sqrt{\sigma_x^2/n + \sigma_y^2/m}}$$

- ii Create two new data sets whose values are $x_i^1 = x_i - \bar{x} + \bar{z}$ and $y_i^1 = y_i - \bar{y} + \bar{z}$ where \bar{z} is the mean of the combined samples
- iii Draw a random sample x_i^* of size n with replacement from x_i^1 and another random sample y_i^* of size m from replacement from y_i^1
- iv Calculate the test statistic

$$t^* = \frac{\bar{x}^* - \bar{y}^*}{\sqrt{\sigma_x^{*2}/n + \sigma_y^{*2}/m}}$$

- v Repeat iii and iv B times to collect B values of the statistic
- vi Estimate the p-value as

$$B = \frac{\sum_{i=1}^B I\{|t_i^*| \geq |t|\}}{B}$$

- (a) First, write an R function that does sampling with replacement for some vector \mathbf{x} . Use `runif` as a source of randomness, and use `ceiling` as a way of rounding continuously distributed numbers to integers.
- (b) What is the computational complexity of the function you just wrote? Assume that `runif` takes n time when requesting n random numbers, and that accessing an element of a vector is free.
- (c) Verify the computational complexity class of your function by running it for variously sized input, and plotting input size versus run time. To get a clearer picture, you might need to average your function call for a given size over multiple runs
- (d) Evaluate the computational complexity of the algorithm at the start of this question. What is its computational complexity?
- (e) Write the above algorithm into R *Optional, just use the provided version of this function in the provided code*
- (f) Check the computational complexity of this implementation matches its theoretical expectations by plotting run time for values of B , m and n . Is it as expected?